

## 2\_广义线性模型

### 2\_1\_线性回归

#### 2\_1\_1\_多元线性回归模型

给定训练数据集

$$D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_i, y_i), \dots, (\mathbf{x}_N, y_N)\}$$

其中,  $\mathbf{x}_i \in \mathcal{X} \subseteq \mathbb{R}^n, y_i \in \mathcal{Y} \subseteq \mathbb{R}$ 。

线性回归模型:

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b = \sum_{i=1}^n w^{(i)} \cdot x^{(i)} + b$$

其中,  $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^n$  是输入,  $\mathbf{w} = (w^{(1)}, w^{(2)}, \dots, w^{(n)})^\top \in \mathbb{R}^n$  和  $b \in \mathbb{R}$  是参数,  $\mathbf{w}$  称为权值向量,  $b$  称为偏置,  $\mathbf{w} \cdot \mathbf{x}$  为  $\mathbf{w}$  和  $\mathbf{x}$  的内积。

令

$$\begin{aligned}\hat{\mathbf{w}} &= (\mathbf{w}, b)^\top \\ \hat{\mathbf{x}} &= (\mathbf{x}, 1)^\top\end{aligned}$$

则多元线性回归模型可简化为

$$f(\hat{\mathbf{x}}) = \hat{\mathbf{w}} \cdot \hat{\mathbf{x}}$$

其中,  $\hat{\mathbf{x}}$  为增广特征向量,  $\hat{\mathbf{w}}$  为增广权重。

#### 2\_1\_2\_多元线性回归参数学习——经验风险最小化与结构风险最小化

损失函数: 平方损失损失函数

$$L(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$$

经验风险

$$R_{emp}(f) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(\mathbf{x}_i))$$

模型参数最优解:

$$\hat{\mathbf{w}}^* = \arg \min_{\hat{\mathbf{w}}} \sum_{i=1}^N (y_i - f(\hat{\mathbf{x}}_i))^2$$

基于均方误差最小化来进行模型求解的方法称为“最小二乘法” (least square method)。

等价的，模型参数最优解：

$$\hat{\mathbf{w}}^* = \arg \min_{\hat{\mathbf{w}}} (\mathbf{y} - \mathbf{X}\hat{\mathbf{w}})^\top (\mathbf{y} - \mathbf{X}\hat{\mathbf{w}})$$

其中，

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top & 1 \\ \mathbf{x}_2^\top & 1 \\ \vdots & \vdots \\ \mathbf{x}_N^\top & 1 \end{pmatrix}$$

$$\mathbf{y} = (y_1, y_2, \dots, y_N)^\top$$

令  $E_{\hat{\mathbf{w}}} = (\mathbf{y} - \mathbf{X}\hat{\mathbf{w}})^\top (\mathbf{y} - \mathbf{X}\hat{\mathbf{w}})$ ，对  $\hat{\mathbf{w}}$  求偏导，得

$$\frac{\partial E_{\hat{\mathbf{w}}}}{\partial \hat{\mathbf{w}}} = 2\mathbf{X}^\top (\mathbf{X}\hat{\mathbf{w}} - \mathbf{y})$$

当  $\mathbf{X}^\top \mathbf{X}$  为满秩矩阵或正定矩阵时，令上式为零可得最优闭式解

$$\hat{\mathbf{w}}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

当上述条件不满足时，可使用主成分分析（PCA）等方法消除特征间的线性相关性，再使用最小二乘法求解。或者通过梯度下降法，初始化  $\hat{\mathbf{w}}_0 = \mathbf{0}$ ，进行迭代

$$\hat{\mathbf{w}} \leftarrow \hat{\mathbf{w}} + \alpha \mathbf{X}^\top (\mathbf{X}\hat{\mathbf{w}} - \mathbf{y})$$

其中， $\alpha$  是学习率。

岭回归（Ridge Regression）正则化项：

$$\alpha \|\mathbf{w}\|^2, \alpha \geq 0.$$

套索回归（Lasso Regression）正则化项：

$$\alpha \|\mathbf{w}\|_1, \alpha \geq 0.$$

弹性网络回归（Elastic Net）正则化项：

$$\alpha \rho \|\mathbf{w}\|_1 + \frac{\alpha(1-\rho)}{2} \|\mathbf{w}\|^2, \alpha \geq 0, 1 \geq \rho \geq 0.$$

### 2\_1\_3\_多元线性回归参数学习——最大似然估计

设标记  $y$  为服从均值为  $f(\hat{\mathbf{x}}) = \hat{\mathbf{w}} \cdot \hat{\mathbf{x}}$ ，方差为  $\sigma^2$  的高斯分布

$$p(y|\hat{\mathbf{x}}, \hat{\mathbf{w}}, \sigma) = \mathcal{N}(y|\hat{\mathbf{w}} \cdot \hat{\mathbf{x}}, \sigma^2)$$

$$= \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y - \hat{\mathbf{w}} \cdot \hat{\mathbf{x}})^2}{2\sigma^2}\right)$$

参数 $\hat{\mathbf{w}}$ 在训练集 $D$ 上的似然函数 (Likelihood)

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}, \hat{\mathbf{w}}, \sigma) &= \prod_{i=1}^N p(y_i|\mathbf{x}_i, \hat{\mathbf{w}}, \sigma) \\ &= \prod_{i=1}^N \mathcal{N}(y_i|\hat{\mathbf{w}} \cdot \hat{\mathbf{x}}_i, \sigma^2) \end{aligned}$$

参数 $\hat{\mathbf{w}}$ 在训练集 $D$ 上的对数似然函数 (Log Likelihood)

$$\log p(\mathbf{y}|\mathbf{X}, \hat{\mathbf{w}}, \sigma) = \sum_{i=1}^N \log \mathcal{N}(y_i|\hat{\mathbf{w}} \cdot \hat{\mathbf{x}}_i, \sigma^2)$$

最大似然估计

$$\hat{\mathbf{w}}^* = \arg \max_{\hat{\mathbf{w}}} p(\mathbf{y}|\mathbf{X}, \hat{\mathbf{w}}, \sigma)$$

等价的

$$\hat{\mathbf{w}}^* = \arg \max_{\hat{\mathbf{w}}} \log p(\mathbf{y}|\mathbf{X}, \hat{\mathbf{w}}, \sigma)$$

等价的

$$\hat{\mathbf{w}}^* = \arg \min_{\hat{\mathbf{w}}} -\log p(\mathbf{y}|\mathbf{X}, \hat{\mathbf{w}}, \sigma)$$

令  $\frac{\partial \log p(\mathbf{y}|\mathbf{X}, \hat{\mathbf{w}}, \sigma)}{\partial \hat{\mathbf{w}}} = 0$ , 得

$$\mathbf{w}^{ML} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

## 2\_1\_4\_多元线性回归模型应用

```
In [3]: import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model, model_selection

def load_data():
    diabetes = datasets.load_diabetes()
    return model_selection.train_test_split(diabetes.data, diabetes.target, test_size=0.25, random_state=0)

def test_LinearRegression(*data):
    X_train, X_test, y_train, y_test = data
    regr = linear_model.LinearRegression()
    regr.fit(X_train, y_train)
    print('Coefficients: %s, intercept %.2f' % (regr.coef_, regr.intercept_))
    print("Residual sum of squares: %.2f" % np.mean((regr.predict(X_test) - y_test) ** 2))
    print('Score: %.2f' % regr.score(X_test, y_test))

if __name__ == '__main__':
    X_train, X_test, y_train, y_test = load_data()
    test_LinearRegression(X_train, X_test, y_train, y_test)
```

```
Coefficients: [ -43.26774487 -208.67053951  593.39797213  302.89814903
-560.27689824
 261.47657106  -8.83343952  135.93715156  703.22658427  28.34844354], intercept 153.07
Residual sum of squares: 3180.20
Score: 0.36
```

```

In [12]: import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model, model_selection

def load_data():
    diabetes = datasets.load_diabetes()
    return model_selection.train_test_split(diabetes.data, diabetes.target,
        test_size=0.25, random_state=0)

def test_Lasso(*data):
    X_train, X_test, y_train, y_test = data
    regr = linear_model.Lasso()
    regr.fit(X_train, y_train)
    print('Coefficients: %s, intercept %.2f' % (regr.coef_, regr.intercept_))
    print("Residual sum of squares: %.2f" % np.mean((regr.predict(X_test) - y_test)
        ** 2))
    print('Score: %.2f' % regr.score(X_test, y_test))

def test_Lasso_alpha(*data):
    X_train, X_test, y_train, y_test = data
    alphas = [0.01, 0.02, 0.05, 0.1, 0.2, 0.5, 1, 2, 5, 10, 20, 50, 100, 200, 500, 1000]
    scores = []
    for i, alpha in enumerate(alphas):
        regr = linear_model.Lasso(alpha=alpha)
        regr.fit(X_train, y_train)
        scores.append(regr.score(X_test, y_test))

    fig = plt.figure()
    ax = fig.add_subplot(1, 1, 1)
    ax.plot(alphas, scores)
    ax.set_xlabel(r"$\alpha$")
    ax.set_ylabel(r"score")
    ax.set_xscale('log')
    ax.set_title("Lasso")
    plt.show()


if __name__ == '__main__':
    X_train, X_test, y_train, y_test = load_data()
    test_Lasso(X_train, X_test, y_train, y_test)
    test_Lasso_alpha(X_train, X_test, y_train, y_test)

```

```

Coefficients: [ 0.         -0.         442.67992538  0.         0.
 0.         -0.         0.         330.76014648  0.         ], intercept
152.52
Residual sum of squares: 3583.42
Score: 0.28

```



```

In [11]: import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model, model_selection

def load_data():
    diabetes = datasets.load_diabetes()
    return model_selection.train_test_split(diabetes.data, diabetes.target,
        test_size=0.25, random_state=0)

def test_Ridge(*data):
    X_train, X_test, y_train, y_test = data
    regr = linear_model.Ridge()
    regr.fit(X_train, y_train)
    print('Coefficients: %s, intercept %.2f' % (regr.coef_, regr.intercept_))
    print("Residual sum of squares: %.2f" % np.mean((regr.predict(X_test) - y_test)
    ** 2))
    print('Score: %.2f' % regr.score(X_test, y_test))

def test_Ridge_alpha(*data):
    X_train, X_test, y_train, y_test = data
    alphas = [0.01, 0.02, 0.05, 0.1, 0.2, 0.5, 1, 2, 5, 10, 20, 50, 100, 200, 500, 1000]
    scores = []
    for i, alpha in enumerate(alphas):
        regr = linear_model.Ridge(alpha=alpha)
        regr.fit(X_train, y_train)
        scores.append(regr.score(X_test, y_test))

    fig = plt.figure()
    ax = fig.add_subplot(1, 1, 1)
    ax.plot(alphas, scores)
    ax.set_xlabel(r"$\alpha$")
    ax.set_ylabel(r"score")
    ax.set_xscale('log')
    ax.set_title("Ridge")
    plt.show()

if __name__ == '__main__':
    X_train, X_test, y_train, y_test = load_data()
    test_Ridge(X_train, X_test, y_train, y_test)
    test_Ridge_alpha(X_train, X_test, y_train, y_test)

```

```

Coefficients: [ 21.19927911 -60.47711393 302.87575204 179.41206395
8.90911449
-28.8080548 -149.30722541 112.67185758 250.53760873 99.57749017], intercept 152.45
Residual sum of squares: 3192.33
Score: 0.36

```

